

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

Why

How?

Who?

Summary

# Cockroach

# A Scalable, Geo-Replicated, Transactional Datastore

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

Why

How?

Who?

Summary

# What kind of datastore?

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

Why

How?

Who?

Summary

- currently a sorted key-value store, but:
- structured and SQL-like layers are coming
- in the end, it should feel like a SQL database (unless you want the lower layers!) with indexes, joins and more
- written in Go

# Scalable, Geo-Replicated, Transactional

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

Why

How?

Who?

Summary

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

Why

How?

Who?

Summary

???

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

Why

How?

Who?

Summary

Cockroach

???

**Scalable**

Geo-  
Replicated

Transactional

Why

How?

Who?

Summary

# Scalable

Cockroach

???

Scalable

**Geo-  
Replicated**

Transactional

Why

How?

Who?

Summary

# Geo-Replicated



Cockroach

???

Scalable

Geo-  
Replicated

**Transactional**

Why

How?

Who?

Summary

# Transactional

# separates Cockroach from NoSQL:

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

Why

How?

Who?

Summary

Consistent & Highly Available is *difficult*:

- apps can do it, but it is very hard (think: encryption)
- the database should do this once, correctly
- the cost is consensus latency
- CockroachDB has transactions that fully deserve the name

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

Why

How?

Who?

Summary

```
opts := client.TransactionOptions{Name: "example put"}
c.RunTransaction(&opts, func(txn *client.KV) error {
    // serializable context!
    gr := proto.GetResponse{}
    txn.Call(proto.Get,
              proto.GetArgs(proto.Key("key1")), &gr)
    txn.Call(proto.Put,
              proto.PutArgs(proto.Key("key2"),
                             append(gr.Value.Bytes, []byte("-new"))),
              &proto.PutResponse{})
    return nil
})
```

Cockroach

???

Scalable

Geo-  
Replicated

**Transactional**

Why

How?

Who?

Summary

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

**Why**

How?

Who?

Summary

# Why

# History

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

Why

How?

Who?

Summary

- SQL “not” scalable or highly available, but transactional
- PostgreSQL, MySQL, Oracle, DB2, ...
- NoSQL scalable and highly available, but not consistent
- BigTable, Cassandra, ...
- NewSQL scalable, highly available, transactions
- Spanner, CockroachDB, ...

# History at Google

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

Why

How?

Who?

Summary

- 2004: BigTable
- 2006: Megastore (on top of BigTable)
- transactional (but slow and complex)
- 2012: Spanner
- fully linearizable (hence consistent)

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

Why

How?

Who?

Summary

*“We believe it is better to have application programmers deal with performance problems due to overuse of transactions as bottlenecks arise, rather than always coding around the lack of transactions.”*

Corbett et al., Spanner paper, 2012



# Google Spanner

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

Why

How?

Who?

Summary

is basically what you would get if SQL and NoSQL had a designer baby that combined both their advantages: scalable, highly available, transactional

but...

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

Why

How?

Who?

Summary

- only Google can have it
- hardware: atomic clocks, GPS receivers
- inhomogeneous infrastructure: TrueTime API, Chubby, Colossus, ...

# We want you

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

Why

How?

Who?

Summary

- to have something like Spanner
- platform semirelational database
- fault-tolerant, transactional, scalable, fast (enough)
- but simpler than Spanner
- simple homogenous infrastructure
- no hardware requirements
- and OpenSource
- this stuff is hard - trust nobody
- see: Jepsen series

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

**Why**

How?

Who?

Summary

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

Why

**How?**

Who?

Summary

# How?

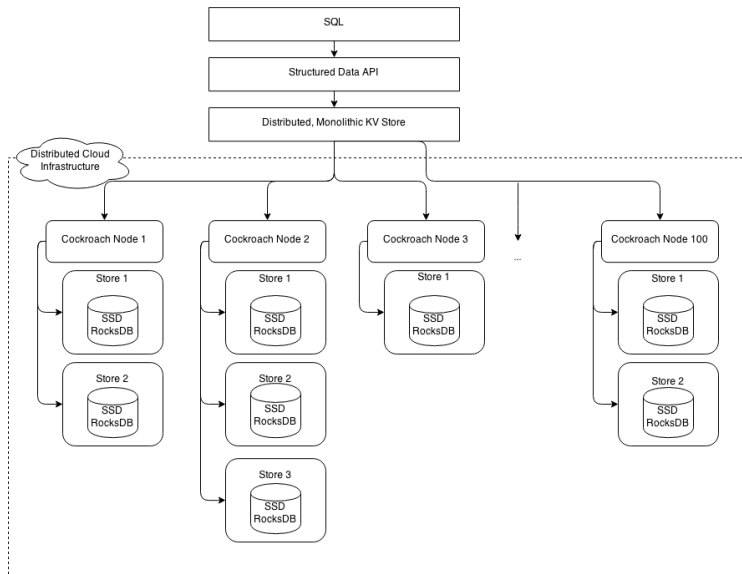


Figure 1: Cockroach Architecture

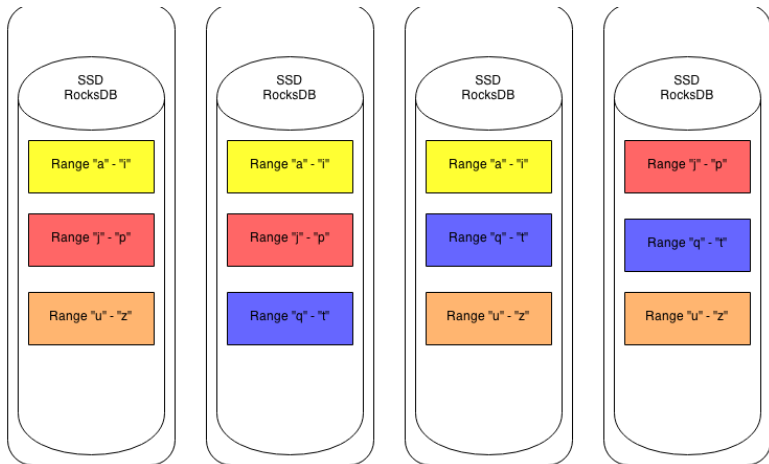


Figure 2: Cockroach Ranges

# Distributed Transactions

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

Why

How?

Who?

Summary

- lock free
- serializable snapshot isolation semantics
- transactions logically don't overlap
- transaction restarts are expected (and normal)
- linearizability for common cases
- a rare concern in practice
- can enforce for all cases when time signal is good



# Under The Hood

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

Why

How?

Who?

Summary

- variation of two phase commit
- txn writes stored as MVCC “intents”
- central transaction table:
  - single key/txn: status, timestamp, priority, ...
  - modified by concurrent txns - first writer wins
  - the single source of truth
  - 2nd phase more efficient – 1 write to transaction table entry
  - intents resolved after commit - correctness doesn't need it!

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

Why

**How?**

Who?

Summary

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

Why

How?

**Who?**

Summary

# Who?

Cockroach

???

Scalable

Geo-  
Replicated

Transactional

Why

How?

Who?

Summary

# Summary